

15Z332 Ex7 - NaiveBayes and SVM

October 17, 2018

1 Exercise 7

1.1 NaiveBayes and SVM Classifiers

Perform classification using NaiveBayes and SVM on Iris dataset. Display the confusion matrix for the various classifiers. Analyze the performance of SVM for different kernel functions. Represent the results of the classification in terms of precision and recall using graphs.

1.1.1 Step 1: Import Iris dataset

```
In [81]: import pandas as pd
         df = pd.read_csv('./Iris.csv')
         df = df.drop('Id',axis=1)
```

1.1.2 Step 2: Convert categorical data into numerical data.

Since the target attribute 'Species' is categorical, we convert it into numerical data using LabelEncoder

```
In [82]: print(df.Species.dtype)
```

object

```
In [83]: from sklearn import preprocessing
```

```
if df.Species.dtype == 'object':
    lbl = preprocessing.LabelEncoder()
    lbl.fit(list(df.Species.values))
    df.Species = lbl.transform(list(df.Species.values))
```

```
print(df.Species.dtype)
```

int64

1.1.3 Step 3: Split the dataset into train and test data

```
In [84]: from sklearn.model_selection import train_test_split
```

```
Y = df.Species.values
X = df.drop(['Species'],axis=1).values
```

```
x_train, x_test, y_train, y_test = train_test_split(X,Y, test_size= 0.20, random_stat
```

1.1.4 Step 4: Fit the data into NaiveBayes classifier

```
In [85]: from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
gnb = GaussianNB()
NaiveBayesModel = gnb.fit(x_train, y_train)
y_pred = NaiveBayesModel.predict(x_test)
accuracy_score(y_test, y_pred)
```

```
Out[85]: 0.9
```

Draw the confusion matrix


```
In [86]: from sklearn.metrics import confusion_matrix
con_mat = confusion_matrix(y_test,y_pred)
print(con_mat)
```

```
[[ 7  0  0]
 [ 0 10  1]
 [ 0  2 10]]
```

```
In [87]: import seaborn as sn
```

```
sn.set(font_scale=1.4)
sn.heatmap(pd.DataFrame(con_mat), annot=True,annot_kws={"size": 16})
```

```
Out[87]: <matplotlib.axes._subplots.AxesSubplot at 0x11a8107f0>
```



1.1.5 Step 5: Fit the data into SVM classifier

In [88]: `from sklearn import svm`

```
from sklearn.metrics import average_precision_score
```

```
x_train, x_test, y_train, y_test = train_test_split(X[Y<2],Y[Y<2], test_size= 0.95)
```

```
classifier = svm.LinearSVC(random_state=27)
```

```
classifier.fit(x_train, y_train)
```

```
y_score = classifier.decision_function(x_test)
```

```
average_precision = average_precision_score(y_test, y_score)
```

```
print('Average precision-recall score: {0:0.2f}'.format(average_precision))
```

Average precision-recall score: 1.00

Since the precision-recall score is 100%, we will try adding noise data and build the SVM model.

In [89]: `import numpy as np`

```
# Add noisy features
```

```
random_state = np.random.RandomState(0)
```

```
n_samples, n_features = X.shape
```

```
X = np.c_[X, random_state.randn(n_samples, 200 * n_features)]

x_train, x_test, y_train, y_test = train_test_split(X[Y<2],Y[Y<2], test_size= 0.20)
```

1) Using Linear kernel

```
In [90]: from sklearn.metrics import precision_recall_curve
import matplotlib.pyplot as plt


classifier = svm.SVC(kernel='linear', random_state=27)
classifier.fit(x_train, y_train)
y_score = classifier.decision_function(x_test)

print('Average precision-recall score:',average_precision_score(y_test, y_score))
precision, recall, _ = precision_recall_curve(y_test, y_score)

plt.step(recall, precision, color='b')
plt.xlabel('Recall')
plt.ylabel('Precision')
```

Average precision-recall score: 0.7725782988940885

Out[90]: Text(0,0.5,'Precision')



2) Using Polynomial kernel


```
In [91]: classifier = svm.SVC(kernel='poly', random_state=27)
classifier.fit(x_train, y_train)
y_score = classifier.decision_function(x_test)

print('Average precision-recall score:', average_precision_score(y_test, y_score))
precision, recall, _ = precision_recall_curve(y_test, y_score)

plt.step(recall, precision, color='r')
plt.xlabel('Recall')
plt.ylabel('Precision')
```

Average precision-recall score: 0.8826354870762766

Out[91]: Text(0,0.5, 'Precision')



3) Using RBF (Radial Basis Function) Kernel

```
In [92]: classifier = svm.SVC(kernel='rbf', random_state=27)
classifier.fit(x_train, y_train)
y_score = classifier.decision_function(x_test)
```

```
print('Average precision-recall score:',average_precision_score(y_test, y_score))
precision, recall, _ = precision_recall_curve(y_test, y_score)

plt.step(recall, precision, color='g')
plt.xlabel('Recall')
plt.ylabel('Precision')
```

Average precision-recall score: 0.7730399230399231

Out[92]: Text(0,0.5, 'Precision')

